

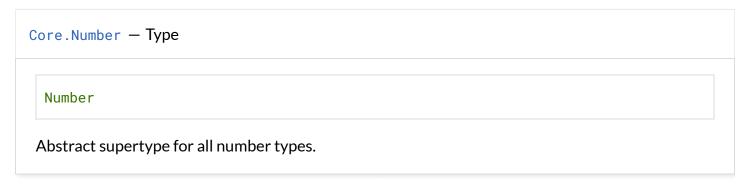




Numbers

Standard Numeric Types

Abstract number types



```
Core.Real - Type
  Real <: Number
Abstract supertype for all real numbers.
```

```
Core.AbstractFloat — Type
  AbstractFloat <: Real
Abstract supertype for all floating point numbers.
```

```
Core.Integer — Type
```

3/20/21, 12:10 1 of 34

```
Integer <: Real</pre>
```

Abstract supertype for all integers.

```
Core.Signed — Type
```

```
Signed <: Integer
```

Abstract supertype for all signed integers.

```
Core.Unsigned — Type
```

```
Unsigned <: Integer
```

Abstract supertype for all unsigned integers.

Base.AbstractIrrational — Type

```
AbstractIrrational <: Real
```

Number type representing an exact irrational value, which is automatically rounded to the correct precision in arithmetic operations with other numeric quantities.

```
Subtypes MyIrrational <: AbstractIrrational should implement at least
==(::MyIrrational, ::MyIrrational), hash(x::MyIrrational, h::UInt), and
convert(::Type{F}, x::MyIrrational) where {F <: Union{BigFloat,Float32,Float64}}.</pre>
```

If a subtype is used to represent values that may occasionally be rational (e.g. a square-root type that represents \sqrt{n} for integers n will give a rational result when n is a perfect square), then it should also implement isinteger, iszero, isone, and == with Real values (since all of these default to false for AbstractIrrational types), as well as defining hash to equal that of the corresponding Rational.

Concrete number types

```
Core.Float16 — Type
```

Float16 <: AbstractFloat

16-bit floating point number type (IEEE 754 standard).

Binary format: 1 sign, 5 exponent, 10 fraction bits.

```
Core.Float32 — Type
```

Float32 <: AbstractFloat

32-bit floating point number type (IEEE 754 standard).

Binary format: 1 sign, 8 exponent, 23 fraction bits.

```
Core.Float64 — Type
```

Float64 <: AbstractFloat

64-bit floating point number type (IEEE 754 standard).

Binary format: 1 sign, 11 exponent, 52 fraction bits.

```
Base.MPFR.BigFloat — Type
```

BigFloat <: AbstractFloat</pre>

Arbitrary precision floating point number type.

Core.Bool — Type

```
Bool <: Integer
```

Boolean type, containing the values true and false.

Bool is a kind of number: false is numerically equal to 0 and true is numerically equal to 1. Moreover, false acts as a multiplicative "strong zero":

```
julia> false == 0
true

julia> true == 1
true

julia> 0 * NaN
NaN

julia> false * NaN
0.0
```

```
Core.Int8 — Type
```

```
Int8 <: Signed</pre>
```

8-bit signed integer type.

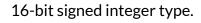
```
Core.UInt8 — Type
```

```
UInt8 <: Unsigned
```

8-bit unsigned integer type.

```
Core.Int16 — Type
```

```
Int16 <: Signed
```



Core.UInt16 — Type

UInt16 <: Unsigned

16-bit unsigned integer type.

Core.Int32 — Type

Int32 <: Signed</pre>

32-bit signed integer type.

Core.UInt32 — Type

UInt32 <: Unsigned

32-bit unsigned integer type.

Core.Int64 — Type

Int64 <: Signed</pre>

64-bit signed integer type.

 ${\tt Core.UInt64-Type}$

UInt64 <: Unsigned

64-bit unsigned integer type.

```
Core.Int128 — Type
```

Int128 <: Signed</pre>

128-bit signed integer type.

Core.UInt128 — Type

UInt128 <: Unsigned

128-bit unsigned integer type.

Base.GMP.BigInt — Type

BigInt <: Signed

Arbitrary precision integer type.

Base.Complex — Type

Complex{T<:Real} <: Number</pre>

Complex number type with real and imaginary part of type T.

ComplexF16, ComplexF32 and ComplexF64 are aliases for Complex{Float16},

Complex{Float32} and Complex{Float64} respectively.

Base.Rational - Type

```
Rational{T<:Integer} <: Real
```

Rational number type, with numerator and denominator of type T. Rationals are checked for overflow.

```
Base.Irrational - Type

Irrational{sym} <: AbstractIrrational</pre>
```

Number type representing an exact irrational value denoted by the symbol sym.

Data Formats

```
Base.digits — Function

digits([T<:Integer], n::Integer; base::T = 10, pad::Integer = 1)</pre>
```

Return an array with element type T (default Int) of the digits of n in the given base, optionally padded with zeros to a specified size. More significant digits are at higher indices, such that $n = sum([digits[k]*base^{(k-1)} for k=1:length(digits)])$.

Examples

```
julia> digits(10, base = 10)
2-element Array{Int64,1}:
0
1

julia> digits(10, base = 2)
4-element Array{Int64,1}:
0
1

julia> digits(10, base = 2, pad = 6)
6-element Array{Int64,1}:
```

```
0
1
0
1
0
0
0
```

```
Base.digits! — Function
```

```
digits!(array, n::Integer; base::Integer = 10)
```

Fills an array of the digits of n in the given base. More significant digits are at higher indices. If the array length is insufficient, the least significant digits are filled up to the array length. If the array length is excessive, the excess portion is filled with zeros.

Examples

```
julia> digits!([2,2,2,2], 10, base = 2)
4-element Array{Int64,1}:
0
1
0
1
julia> digits!([2,2,2,2,2,2], 10, base = 2)
6-element Array{Int64,1}:
0
1
0
1
0
1
0
1
0
1
```

```
Base.bitstring — Function
```

```
bitstring(n)
```

A string giving the literal bit representation of a number.

Examples

Base.parse — Function

```
parse(type, str; base)
```

Parse a string as a number. For Integer types, a base can be specified (the default is 10). For floating-point types, the string is parsed as a decimal floating-point number. Complex types are parsed from decimal strings of the form "R±Iim" as a Complex(R,I) of the requested type; "i" or "j" can also be used instead of "im", and "R" or "Iim" are also permitted. If the string does not contain a valid number, an error is raised.

```
Julia 1.1
parse(Bool, str) requires at least Julia 1.1.
```

Examples

```
julia> parse(Int, "1234")
1234

julia> parse(Int, "1234", base = 5)
194

julia> parse(Int, "afc", base = 16)
2812

julia> parse(Float64, "1.2e-3")
0.0012

julia> parse(Complex{Float64}, "3.2e-1 + 4.5im")
0.32 + 4.5im
```

```
Base.tryparse — Function
```

```
tryparse(type, str; base)
```

Like parse, but returns either a value of the requested type, or nothing if the string does not contain a valid number.

```
Base.big — Function
```

```
big(x)
```

Convert a number to a maximum precision representation (typically BigInt or BigFloat). See BigFloat for information about some pitfalls with floating-point numbers.

```
Base.signed — Function
```

```
signed(T::Integer)
```

Convert an integer bitstype to the signed type of the same size.

Examples

```
julia> signed(UInt16)
Int16
julia> signed(UInt64)
Int64
```

```
signed(x)
```

Convert a number to a signed integer. If the argument is unsigned, it is reinterpreted as signed without checking for overflow.

Base.unsigned — Function

```
unsigned(T::Integer)
```

Convert an integer bitstype to the unsigned type of the same size.

Examples

```
julia> unsigned(Int16)
UInt16
```

```
julia> unsigned(UInt64)
UInt64
```

```
Base.float - Method
```

```
float(x)
```

Convert a number or array to a floating point data type.

```
Base.Math.significand — Function
```

```
significand(x)
```

Extract the significand(s) (a.k.a. mantissa), in binary representation, of a floating-point number. If x is a non-zero finite number, then the result will be a number of the same type on the interval [1,2). Otherwise x is returned.

Examples

```
julia> significand(15.2)/15.2
0.125

julia> significand(15.2)*8
15.2
```

Base.Math.exponent — Function

```
exponent(x) -> Int
```

Get the exponent of a normalized floating-point number.

```
Base.complex — Method
```

```
complex(r, [i])
```

Convert real numbers or arrays to complex. i defaults to zero.

Examples

```
julia> complex(7)
7 + 0im

julia> complex([1, 2, 3])
3-element Array{Complex{Int64},1}:
1 + 0im
2 + 0im
3 + 0im
```

Base.bswap — Function

```
bswap(n)
```

Reverse the byte order of n.

(See also ntoh and hton to convert between the current native byte order and big-endian order.)

Examples

```
Base.hex2bytes — Function
```

```
hex2bytes(s::Union{AbstractString, AbstractVector{UInt8}})
```

Given a string or array s of ASCII codes for a sequence of hexadecimal digits, returns a Vector {UInt8} of bytes corresponding to the binary representation: each successive pair of hexadecimal digits in s gives the value of one byte in the return vector.

The length of s must be even, and the returned array has half of the length of s. See also hex2bytes! for an in-place version, and bytes2hex for the inverse.

Examples

```
julia> s = string(12345, base = 16)
"3039"
julia> hex2bytes(s)
2-element Array{UInt8,1}:
 0x30
 0x39
julia> a = b"01abEF"
6-element Base.CodeUnits{UInt8,String}:
 0x30
 0x31
 0x61
 0x62
 0x45
 0x46
julia> hex2bytes(a)
3-element Array{UInt8,1}:
 0x01
 0xab
 0xef
```

```
Base.hex2bytes! — Function
```

```
hex2bytes!(d::AbstractVector{UInt8}, s::Union{String,AbstractVector{UInt8}})
```

Convert an array s of bytes representing a hexadecimal string to its binary representation, similar to hex2bytes except that the output is written in-place in d. The length of s must be exactly twice

the length of d.

Base.bytes2hex — Function

```
bytes2hex(a::AbstractArray{UInt8}) -> String
bytes2hex(io::I0, a::AbstractArray{UInt8})
```

Convert an array a of bytes to its hexadecimal string representation, either returning a String via bytes2hex(a) or writing the string to an io stream via bytes2hex(io, a). The hexadecimal characters are all lowercase.

Examples

```
julia> a = string(12345, base = 16)
"3039"

julia> b = hex2bytes(a)
2-element Array{UInt8,1}:
    0x30
    0x39

julia> bytes2hex(b)
"3039"
```

General Number Functions and Constants

```
Base.one — Function
```

```
one(x)
one(T::type)
```

Return a multiplicative identity for x: a value such that one(x)*x == x*one(x) == x. Alternatively one(T) can take a type T, in which case one returns a multiplicative identity for any x of type T.

If possible, one(x) returns a value of the same type as x, and one(T) returns a value of type T. However, this may not be the case for types representing dimensionful quantities (e.g. time in

days), since the multiplicative identity must be dimensionless. In that case, one(x) should return an identity value of the same precision (and shape, for matrices) as x.

If you want a quantity that is of the same type as x, or of type T, even if x is dimensionful, use one unit instead.

Examples

```
julia> one(3.7)
1.0

julia> one(Int)
1

julia> import Dates; one(Dates.Day(1))
1
```

Base.oneunit — Function

```
oneunit(x::T)
oneunit(T::Type)
```

Returns T(one(x)), where T is either the type of the argument or (if a type is passed) the argument. This differs from one for dimensionful quantities: one is dimensionless (a multiplicative identity) while oneunit is dimensionful (of the same type as x, or of type T).

Examples

```
julia> oneunit(3.7)
1.0

julia> import Dates; oneunit(Dates.Day)
1 day
```

```
Base.zero - Function
```

```
zero(x)
```

Get the additive identity element for the type of x (x can also specify the type itself).

Examples

```
julia> zero(1)
0

julia> zero(big"2.0")
0.0

julia> zero(rand(2,2))
2×2 Array{Float64,2}:
0.0 0.0
0.0 0.0
```

Base.im — Constant

im

The imaginary unit.

Examples

```
julia> im * im
-1 + 0im
```

Base.MathConstants.pi — Constant

```
\boldsymbol{\pi} pi
```

The constant pi.

Examples

```
julia> pi
π = 3.1415926535897...
```

```
Base.MathConstants.e — Constant
The constant e.
Examples
  julia> e
 e = 2.7182818284590...
Base.MathConstants.catalan — Constant
  catalan
Catalan's constant.
Examples
  julia> Base.MathConstants.catalan
  catalan = 0.9159655941772...
{\tt Base.MathConstants.eulergamma-Constant}
  eulergamma
 Euler's constant.
Examples
  julia> Base.MathConstants.eulergamma
  y = 0.5772156649015...
```

```
Base.MathConstants.golden — Constant
  golden
The golden ratio.
Examples
  julia> Base.MathConstants.golden
  \varphi = 1.6180339887498...
Base.Inf — Constant
  Inf, Inf64
Positive infinity of type Float64.
Base.Inf32 — Constant
  Inf32
Positive infinity of type Float32.
Base.Inf16 — Constant
  Inf16
Positive infinity of type Float16.
Base.NaN — Constant
```

```
NaN, NaN64
```

A not-a-number value of type Float64.

Base.NaN32 — Constant

NaN32

A not-a-number value of type Float32.

Base.NaN16 — Constant

NaN16

A not-a-number value of type Float16.

Base.issubnormal — Function

```
issubnormal(f) -> Bool
```

Test whether a floating point number is subnormal.

Base.isfinite — Function

```
isfinite(f) -> Bool
```

Test whether a number is finite.

Examples

```
julia> isfinite(5)
true
```

```
julia> isfinite(NaN32)
false
```

```
Base.isinf — Function
```

```
isinf(f) -> Bool
```

Test whether a number is infinite.

```
Base.isnan — Function
```

```
isnan(f) -> Bool
```

Test whether a number value is a NaN, an indeterminate value which is neither an infinity nor a finite number ("not a number").

```
Base.iszero — Function
```

```
iszero(x)
```

Return true if x = zero(x); if x is an array, this checks whether all of the elements of x are zero.

Examples

```
julia> iszero(0.0)
true

julia> iszero([1, 9, 0])
false

julia> iszero([false, 0, 0])
true
```

```
Base.isone — Function
```

```
isone(x)
```

Return true if x == one(x); if x is an array, this checks whether x is an identity matrix.

Examples

```
julia> isone(1.0)
true

julia> isone([1 0; 0 2])
false

julia> isone([1 0; 0 true])
true
```

Base.nextfloat — Function

```
nextfloat(x::AbstractFloat, n::Integer)
```

The result of n iterative applications of nextfloat to x if $n \ge 0$, or -n applications of prevfloat if n < 0.

```
nextfloat(x::AbstractFloat)
```

Return the smallest floating point number y of the same type as x such x < y. If no such y exists (e.g. if x is Inf or NaN), then return x.

```
Base.prevfloat — Function
```

```
prevfloat(x::AbstractFloat, n::Integer)
```

The result of n iterative applications of prevfloat to x if n >= 0, or -n applications of nextfloat if n < 0.

```
prevfloat(x::AbstractFloat)
```

Return the largest floating point number y of the same type as x such y < x. If no such y exists (e.g. if x is -Inf or NaN), then return x.

```
Base.isinteger — Function
```

```
isinteger(x) -> Bool
```

Test whether x is numerically equal to some integer.

Examples

```
julia> isinteger(4.0)
true
```

```
Base.isreal — Function
```

```
isreal(x) -> Bool
```

Test whether x or all its elements are numerically equal to some real number including infinities and NaNs. isreal(x) is true if isequal(x, real(x)) is true.

Examples

```
julia> isreal(5.)
true

julia> isreal(Inf + 0im)
true
```

```
julia> isreal([4.; complex(0,1)])
false
```

```
Core.Float32 - Method
```

```
Float32(x [, mode::RoundingMode])
```

Create a Float32 from x. If x is not exactly representable then mode determines how x is rounded.

Examples

```
julia> Float32(1/3, RoundDown)
0.3333333f0

julia> Float32(1/3, RoundUp)
0.33333334f0
```

See RoundingMode for available rounding modes.

Core.Float64 - Method

```
Float64(x [, mode::RoundingMode])
```

Create a Float64 from x. If x is not exactly representable then mode determines how x is rounded.

Examples

```
julia> Float64(pi, RoundDown)
3.141592653589793

julia> Float64(pi, RoundUp)
3.1415926535897936
```

See RoundingMode for available rounding modes.

```
Base.Rounding.rounding — Function
```

```
rounding(T)
```

Get the current floating point rounding mode for type T, controlling the rounding of basic arithmetic functions (+, -, *, / and sqrt) and type conversion.

See RoundingMode for available modes.

Base.Rounding.setrounding — Method

```
setrounding(T, mode)
```

Set the rounding mode of floating point type T, controlling the rounding of basic arithmetic functions (+, -, *, / and sqrt) and type conversion. Other numerical functions may give incorrect or invalid values when using rounding modes other than the default RoundNearest.

Note that this is currently only supported for T == BigFloat.



Warning

This function is not thread-safe. It will affect code running on all threads, but its behavior is undefined if called concurrently with computations that use the setting.

Base.Rounding.setrounding — Method

```
setrounding(f::Function, T, mode)
```

Change the rounding mode of floating point type T for the duration of f. It is logically equivalent to:

```
old = rounding(T)
setrounding(T, mode)
f()
setrounding(T, old)
```

See RoundingMode for available rounding modes.

Base.Rounding.get_zero_subnormals — Function

get_zero_subnormals() -> Bool

Return false if operations on subnormal floating-point values ("denormals") obey rules for IEEE arithmetic, and true if they might be converted to zeros.

• Warning

This function only affects the current thread.

 ${\tt Base.Rounding.set_zero_subnormals} - {\tt Function}$

set_zero_subnormals(yes::Bool) -> Bool

If yes is false, subsequent floating-point operations follow rules for IEEE arithmetic on subnormal values ("denormals"). Otherwise, floating-point operations are permitted (but not required) to convert subnormal inputs or outputs to zero. Returns true unless yes==true but the hardware does not support zeroing of subnormal numbers.

 $set_zero_subnormals(true)$ can speed up some computations on some hardware. However, it can break identities such as (x-y==0) == (x==y).

Warning

This function only affects the current thread.

Integers

Base.count_ones — Function

```
count_ones(x::Integer) -> Integer
```

Number of ones in the binary representation of x.

Examples

```
julia> count_ones(7)
3
```

```
Base.count_zeros - Function
```

```
count_zeros(x::Integer) -> Integer
```

Number of zeros in the binary representation of x.

Examples

```
julia> count_zeros(Int32(2 ^ 16 - 1))
16
```

```
Base.leading_zeros — Function
```

```
leading_zeros(x::Integer) -> Integer
```

Number of zeros leading the binary representation of x.

Examples

```
julia> leading_zeros(Int32(1))
31
```

```
Base.leading_ones — Function
```

```
leading_ones(x::Integer) -> Integer
```

Number of ones leading the binary representation of x.

Examples

```
julia> leading_ones(UInt32(2 ^ 32 - 2))
31
```

```
Base.trailing_zeros - Function
```

```
trailing_zeros(x::Integer) -> Integer
```

Number of zeros trailing the binary representation of x.

Examples

```
julia> trailing_zeros(2)
1
```

Base.trailing_ones — Function

```
trailing_ones(x::Integer) -> Integer
```

Number of ones trailing the binary representation of x.

Examples

```
julia> trailing_ones(3)
2
```

Base.isodd — Function

```
isodd(x::Integer) -> Bool
```

Return true if x is odd (that is, not divisible by 2), and false otherwise.

Examples

```
julia> isodd(9)
true

julia> isodd(10)
false
```

Base.iseven — Function

```
iseven(x::Integer) -> Bool
```

Return true if x is even (that is, divisible by 2), and false otherwise.

Examples

```
julia> iseven(9)
false

julia> iseven(10)
true
```

Core.@int128_str — Macro

```
@int128_str str
@int128_str(str)
```

@int128_str parses a string into a Int128 Throws an ArgumentError if the string is not a valid integer

```
Core.@uint128_str — Macro
```

```
@uint128_str str
@uint128_str(str)
```

@uint128_str parses a string into a UInt128 Throws an ArgumentError if the string is not a valid integer

BigFloats and BigInts

The BigFloat and BigInt types implements arbitrary-precision floating point and integer arithmetic, respectively. For BigFloat the GNU MPFR library is used, and for BigInt the GNU Multiple Precision Arithmetic Library (GMP) is used.

```
Base.MPFR.BigFloat - Method
```

```
\label{eq:bigFloat} BigFloat(x::Union\{Real, AbstractString\} \ [\ , \ rounding::RoundingMode=rounding(BigFloat(x::Union\{Real, AbstractString\}), \ ] \\
```

Create an arbitrary precision floating point number from x, with precision precision. The rounding argument specifies the direction in which the result should be rounded if the conversion cannot be done exactly. If not provided, these are set by the current global values.

BigFloat(x::Real) is the same as convert(BigFloat,x), except if x itself is already BigFloat, in which case it will return a value with the precision set to the current global precision; convert will always return x.

BigFloat(x::AbstractString) is identical to parse. This is provided for convenience since decimal literals are converted to Float64 when parsed, so BigFloat(2.1) may not yield what you expect.



precision as a keyword argument requires at least Julia 1.1. In Julia 1.0 precision is the second positional argument (BigFloat(x, precision)).

Examples

```
julia> BigFloat(2.1) # 2.1 here is a Float64
2.100000000000000088817841970012523233890533447265625
```

See also

- @big_str
- rounding and setrounding
- precision and setprecision

```
Base.precision — Function
```

```
precision(num::AbstractFloat)
```

Get the precision of a floating point number, as defined by the effective number of bits in the significand.

```
Base.precision — Method
```

```
precision(BigFloat)
```

Get the precision (in bits) currently used for BigFloat arithmetic.

```
Base.MPFR.setprecision — Function
```

```
setprecision([T=BigFloat,] precision::Int)
```

Set the precision (in bits) to be used for T arithmetic.



This function is not thread-safe. It will affect code running on all threads, but its behavior is undefined if called concurrently with computations that use the setting.

```
setprecision(f::Function, [T=BigFloat,] precision::Integer)
```

Change the T arithmetic precision (in bits) for the duration of f. It is logically equivalent to:

```
old = precision(BigFloat)
setprecision(BigFloat, precision)
f()
setprecision(BigFloat, old)
```

Often used as setprecision(T, precision) do ... end

Note: nextfloat(), prevfloat() do not use the precision mentioned by setprecision

Base.GMP.BigInt - Method

```
BigInt(x)
```

Create an arbitrary precision integer. x may be an Int (or anything that can be converted to an Int). The usual mathematical operators are defined for this type, and results are promoted to a BigInt.

Instances can be constructed from strings via parse, or using the big string literal.

Examples

Core.@big_str — Macro

```
@big_str str
@big_str(str)
```

Parse a string into a BigInt or BigFloat, and throw an ArgumentError if the string is not a valid number. For integers _ is allowed in the string as a separator.

Examples

```
julia> big"123_456"
123456

julia> big"7891.5"
7891.5
```

« Mathematics Strings »

Powered by Documenter.jl and the Julia Programming Language.